

Wenup

ENGINEERING CANDIDATE TECHNICAL TEST LLM APPLICATION

DOCUMENT INTAKE ASSISTANT

SEPTEMBER 2026



OVERVIEW

THE BRIEF

Build a small web application called Document Intake Assistant. The application should conduct a conversational interview, maintain structured information collected during that conversation, and produce a draft document from the information supplied.

The objective is not to build a sophisticated legal product. This exercise evaluates how you build reliable applications around LLMs, handle state and ambiguity, and combine conventional software engineering with generative AI.

You may use any LLM provider and development framework. Use of AI development tools such as ChatGPT, Claude, Cursor or Copilot is encouraged.

The solution does not need to be production-ready. We are more interested in your engineering decisions and your ability to explain them.

SCENARIO

FICTIONAL PRODUCT

The application is creating a fictional Personal Wishes Document. It must collect the following information through conversation:

- Full name
- Home address
- Whether the document covers worldwide assets
- Whether the user has children
- Names of children, if applicable
- Name of executor
- Executor's relationship to the user
- Any specific gifts
- Any additional wishes

EXAMPLE

Assistant: Who would you like to appoint as your executor?

User: My brother James.

REQUIREMENTS

CORE EXPERIENCE

Application

- Provide both a backend and a simple user interface.
- Let the user have a multi-turn conversation with the assistant.
- Show a live preview of the structured information and the draft document.
- Allow the user to correct previously supplied information.

Structured state

Define an explicit schema for the information collected. Conversation history alone is not sufficient as the application's source of truth.

STRUCTURED STATE

```
{
  "full_name": "Jane Smith",
  "covers_worldwide_assets": true,
  "has_children": false,
  "executor": {
    "name": "James Smith",
    "relationship": "brother"
  }
}
```

LLM BEHAVIOUR

RELIABILITY MATTERS

Conversation and ambiguity

- Ask sensible follow-up questions when an answer is missing, unclear or contradictory.
- Do not invent facts. Represent unknown or unconfirmed values explicitly.
- Avoid repeatedly asking for information that has already been captured.
- Handle answers that provide several fields at once, in any reasonable order.

State and document generation

- Validate model output before applying it to the structured state.
- Keep the preview consistent with the latest confirmed state.
- Generate a clear draft Personal Wishes Document from that state.
- Clearly label the document as fictional and not legal advice.

ENGINEERING

WHAT TO DEMONSTRATE

Technical expectations

- A clear separation between UI, application logic, LLM interaction and document generation.
- A defined API contract and structured data model.
- Graceful handling of model errors, malformed responses and missing configuration.
- A small but meaningful set of automated tests around the most important behaviour.
- Straightforward local setup with secrets kept out of source control.

Use of AI tools

Include a brief AI log with key prompts, notable iterations, and examples of output you questioned or corrected. Copy/paste is fine; concise and candid is better than polished.

We value sound judgement more than the number of features completed.

SUBMISSION

WHAT TO SEND

Submission instructions


- Share a public Git repository containing your solution.
- Include complete instructions for configuring and running the application.
- Include your AI log and a short note on what you would improve for production.

If you cannot access an LLM API

You may use a deterministic mock or local stub behind a clearly defined LLM interface. Your solution should still demonstrate:

- Structured request and response contracts for the model interaction.
- Fixtures covering valid, ambiguous and malformed model responses.
- State updates, corrections, validation, error handling and document generation.
- A short explanation of how the mock would be replaced by a real provider.

Access to a paid LLM API is not required to complete the exercise.



Wenup

THANK YOU

We look forward to seeing how you approach the problem.

WENUP.CO.UK

